




Average Lap Duration is higher-is-better — your example (00:25 vs 00:15 = Improved) confirms it: longer uninterrupted stretches on a task mean less fragmentation. It isn't in either direction list in §8, so it needs adding explicitly, or the trend engine will read it backwards.

Appreciation counts shouldn't go through the trend engine. They're RO-given during weekly evaluation, so real values are 0–3. The FRD's own example shows 1 vs 3 = "Concern" — one fewer appreciation in a quarter labelled as continuous degradation. That's the engine misfiring on small numbers, not a performance signal. I'd show both as period counts with the raw previous value and no qualitative status, which §8 already permits for metrics that don't map cleanly. Same caution for Negative Leaves at 2 vs 0.

Indirect reportees need rater attribution. If a skip-level RO sees 3.2/5 on Customer Governance, the first question is who rated it and in which weeks. I'd keep the main card clean and put the rater name in the 12-week competency drill-down, where the weekly scores already live.

On keeping both surfaces: the division that holds is **Weekly History = the ritual trail** (per-review, immutable, remarks and actions, direct reports only) and **Performance Card = the analysis layer** (any period, self or any reportee, no review content). So the Performance Card's competency drill-down shows scores over 12 weeks — never remarks or actions. That keeps them from reading as two versions of the same screen.

Placement: it defaults to Self and every user has one, so it belongs in the personal/workspace group, not under My Team. The resource dropdown makes the RO case the same screen.

Four calls worth flagging:

Trend word now appears in the summary cards too. Billable Split reads "↓ Concern" up top and shows as Concern in Needs Attention — one vocabulary across the page. The arrow carries direction, the word carries the 3-period judgement.

Work & Productivity and Contribution sit side by side. FRD order is preserved reading left-to-right, but stacking them would have pushed Snapshot far below the fold for no benefit — both are compact tables and neither needs full width.


"Count only" chips on Negative Leaves and both Appreciation rows, with a footnote explaining why. That's the small-N problem handled visibly rather than silently.

Missing data is rendered, not hidden. Tooling & Automation shows a dashed empty bar and "—", with the caption confirming it's excluded from the 3.8 average — so nobody reads a blank as a zero.

Performance Card — 12-Week Metric History (Modal)

there's a real mismatch to solve inside it: the drill-down is **weekly (12 weeks)** while the trend badge is computed on **3 comparable periods** (months, if the page is set to This Month). If the modal only shows weeks, the "Improved" badge floating at the top has no visible basis.

So the modal carries both: weeks as the chart, and the period basis as one explicit strip.



Four calls in there:

The Trend Basis strip is the important one. It shows the three monthly values the badge actually used, and reconciles why the strip says Improving while the card says Improved — March is a partial period. Without this, weekly bars and a period-level badge look like a bug.

Month bands under the week axis tie the 12 weeks back to the periods, so the chart and the basis strip read as one thing rather than two datasets.


Bar shading carries recency, not judgement — the selected period is darkest, the current week darkest of all. I deliberately didn't colour bars good/bad; the four-status vocabulary already owns that, and per-bar colouring would imply weekly thresholds you don't have.

Prev/next metric stepper in the header (↑ Focus Hrs / ↓ Avg Lap Duration) lets an RO walk the table without closing and reopening. That's an addition beyond the FRD — say if you'd rather keep it strictly one-metric-per-open.

The four stat tiles (current, average, range, coverage) are the "no separate report required" line honoured — enough to read the history without a second screen.

Performance Card — 12-Week Competency Detail (Modal)

one modal for both — Behavioural and Technical Competency. And the comparative payoff only lands if each section carries a **per-week average row**, so you can see the two trajectories move against each other rather than eyeballing ten rows.



Four calls:

Tinted score cells, not bars per row. Ten rows × 12 weeks needs a heatmap read — you scan for pale/warm patches, then read the number. Bars would have made it 120 tiny charts.

The per-week average row is where the comparison lives. Behavioural climbing while Technical sits flat at 3.5–3.8 is the finding, and it's visible in seconds by comparing the two strips. That's the argument for one modal rather than two.

Movement column per competency uses the same four-status vocabulary as everywhere else, so Customer Governance reading Concern here matches Needs Attention on the card.

Two distinct kinds of missing data are shown differently in the footnote: W08 (no review that week, whole column dashed) versus Tooling & Automation (competency added mid-quarter, never rated). Both render as —, but the reasons differ and an RO will ask.

The "opened from Behavioural" strip is the only nod to which link was clicked — no scroll jump, since the whole point is seeing both.