

Config layer that everything upstream reads from. Two objects live here, so a segmented page: **Department Competencies** (what the RO rates against) and **Ranking Master** (the scale they rate on). Both admin-only under Setup, mirroring your other masters.

First the competency editor — department-scoped, two capped lists with sequencing and cross-department reuse:

The screenshot shows the 'Weekly Check-In Configuration' interface, specifically the 'Department Competencies' section. The page is titled 'Weekly Check-In Configuration' with a subtitle 'Setup · competency & ranking setup · Super Admin only'. A 'Super Admin' button is in the top right. The left sidebar shows a list of departments: 'Quality Eng.' (checked), 'Platform Eng.' (checked), 'Data & Analytics' (checked), 'Client Delivery' (not set), and 'Design' (not set). The main content area is for 'Quality Engineering' and shows two columns: 'TECHNICAL' (4 of 5) and 'BEHAVIOURAL' (3 of 5). Each column has a list of competencies with drag handles and delete icons. The 'TECHNICAL' column lists: 1. Test Coverage & Rigor, 2. Defect Detection, 3. Automation Practices, and 4. Root-Cause Analysis. The 'BEHAVIOURAL' column lists: 1. Ownership & Reliability, 2. Collaboration, and 3. Communication. Below each list is a button to '+ Add technical/behavioural competency' with the number of slots left. A note at the bottom states: 'Fewer than five is fine — a department only rates the competencies it defines. Empty categories are simply skipped in the workspace.' At the bottom of the page, there is a status bar showing 'Unsaved · added "Root-Cause Analysis"', a 'Discard' button, and a 'Save competencies' button.

Weekly Check-In Configuration
Setup · competency & ranking setup · Super Admin only

Super Admin

Department Competencies

Ranking Master

DEPARTMENTS

- Quality Eng. ✓
- Platform Eng. ✓
- Data & Analytics ✓
- Client Delivery •
- Design not set

Quality Engineering
Competencies apply to all resources in this department

Copy from department...

TECHNICAL · 4 of 5 drag to reorder

- 1 Test Coverage & Rigor
- 2 Defect Detection
- 3 Automation Practices
- 4 Root-Cause Analysis

+ Add technical competency · 1 slot left

BEHAVIOURAL · 3 of 5 drag to reorder

- 1 Ownership & Reliability
- 2 Collaboration
- 3 Communication

+ Add behavioural competency · 2 slots left

Fewer than five is fine — a department only rates the competencies it defines. Empty categories are simply skipped in the workspace.

Unsaved · added "Root-Cause Analysis" Discard Save competencies

And the **Ranking Master** — the second segment. One shared 5-level scale, ordered high→low, titles editable:

Ranking Master

The 5-level scale every competency is rated on - shared across all departments

Reusable by future Performance Intelligence

VALUE	RANKING TITLE · editable	ORDER
5	Outstanding	Highest
4	Exceeds Expectations	▲ ▼
3	Meets Expectations	▲ ▼
2	Needs Improvement	▲ ▼
1	Unsatisfactory	Lowest

Exactly five levels — titles can be renamed, but the high→low order and 5-value scale are fixed so historical ratings stay comparable.

Admin Config — spec

Segmented page (Competencies · Ranking Master) — one Setup destination for the two config objects the RO surfaces depend on, admin-only.

Department Competencies (WCI001–005)


- **Department rail** on the left with a config-status glance: ✓ set, amber partial, red "not set." That last state matters — it's the **exception case** where the workspace *blocks submission* until an admin configures competencies. Surfacing it here is where the admin fixes it before an RO ever hits the wall.
- **Two capped lists**, Technical and Behavioural, each **≤5 (BR001)**. The cap is shown live ("4 of 5", "1 slot left") and the Add affordance disappears at five rather than erroring.
- **Sequencing (WCI004)** via drag handles + the number badge — the order here is the order competencies appear in the workspace.
- **Fewer-than-five is legitimate (WCI025)** — called out explicitly so an admin doesn't feel obligated to fill all ten slots; empty categories just skip in the review.
- **Template reuse (WCI005)** — "Copy from department..." seeds a new department from an existing one, then edit. This is the accelerator for standing up a new department fast.

Ranking Master (WCI006–010)

- **Exactly five levels (WCI007)**, ordered **high→low and locked in order (WCI009)** — you can rename titles (WCI008, the 🖋️) but not add/remove levels or reorder the scale, so ratings stay comparable over time.
- **Value + colour are fixed**, title is the editable part — the scale reads as one object, not five free rows.

- Flagged as **reusable by future PIE (WCI010)** — the small note ties it to the performance-intelligence roadmap the FRD keeps pointing at.

The RO's Monday-morning surface — pick the week, see your direct reports, knock through the pending ones. This is what feeds into the workspace: each "Open review" lands on the two-pane screen I just sketched.



Review Queue — spec

The weekly frame

- Week picker** front and centre (WCI011) with prev/next arrows + a "This week" reset — the RO's first act is choosing which week they're reviewing.
- Header states the boundary:** "your direct reports only" — and the footer repeats it. This is the deliberate **direct-only** line (BR002/WCI012), the contrast with Access Rights' recursive visibility I flagged last turn.

The count strip (WCI013/014)

- Four tiles: Direct reports (5), **Pending (2, amber)**, Completed (3, green), and a **week-progress bar (3 of 5)**. The pending count is the number the RO is actually chasing — it's given its own colour so it reads at a glance.

The queue itself

- Sorted pending-first** so the work-to-do floats up; completed rows dim back. Filter tabs (All / Pending / Completed) + search across your reports.

- Each row previews **why this person needs attention before you open them**: last week's status, this week's Confirmation Discipline (the leading indicator), and any **open action item** carried forward — so the RO triages without opening every review.
- Two row states**: Pending → primary **"Open review"** (lands on the workspace); Completed → quiet **"View"** (view-only, since submitted reviews are immutable per BR008), with the submission day shown.

Edge cases baked in

- Dev Malhotra** — joined mid-week: "No prior review" + "No data" chip, but still reviewable (the exception table allows a review with no operational data).
- The **completed rows carry their outcome** — Vikram's recognition ✓, Ananya's intervention action scheduled — so the queue doubles as a week-at-a-glance record, not just a to-do list.

single-screen two-pane workspace, under **My Team**, Check-In Workspace first.

Left pane is read-only evidence (the RPR/PER metrics, pulled automatically per WCI-BR004); right pane is where the RO does the work — rate competencies, set status/confidence, log the action, close out last week. Header carries the week + which report you're on.

R

< Queue

PN

Priya Nair · QA Engineer · Quality Eng.

Direct report · reviewing as Arjun Mehta

Week of Jan 27

Pending · 2 of 5 in queue

EVIDENCE · FROM ONEVIEW

read-only · auto-pulled

Planning Accuracy

95%

2 deviations / 40 planned

Confirm. Discipline

79%

2 late confirmations

Utilization

32 h

of 40 h capacity

Billable split

66% billable · 34% non

PROJECTS THIS WEEK

Project Atlas

22 h

Project Nova

10 h

LAST WEEK'S REVIEW

Jan 20

Watch

Confidence: Medium

"Confirmation timeliness slipping — asked Priya to confirm daily rather than batching Fridays. Skills solid, delivery on Atlas strong."

Open action · Coaching

Daily confirmation habit

Still pending

Completed

ASSESSMENT & SUBMISSION

TECHNICAL COMPETENCIES

Quality Eng. template · rate 1-5

Test Coverage & Rigor

1 2 3 4 5

Defect Detection

1 2 3 4 5

Automation Practices

1 2 3 4 5

BEHAVIOURAL COMPETENCIES

Ownership & Reliability

1 2 3 4 5

Collaboration

1 2 3 4 5

WEEKLY STATUS

On Track

Watch

Intervention

CONFIDENCE

High

Med

Low

RO REMARKS

min 100 chars · 142 ✓

Strong quarter on Atlas — defect detection is a genuine strength. Confirmation discipline is the one watch item; the daily-confirm habit from last week helped but isn't sticky yet. Keeping at Watch one more week.

ACTION TYPE

Coaching

ACTION NOTES · required · min 100

Pair with senior QA on confirmation workflow; check daily for two weeks then reassess.

RECOGNITION

None

Appreciate

Publicly

Immutable once submitted · one check-in per week

Save draft

Submit review

Weekly Check-In Workspace — spec

The two-pane split

- **Left = evidence, read-only** (WCI016–021, BR004): the same metrics from RPR/PER — Planning Accuracy + deviation count, Confirmation Discipline + late count, Utilization vs capacity, billable split, projects worked. The lock badge + "auto-pulled from OneView" makes it unmistakable this side isn't editable. It's *coaching evidence*, not input.
- **Right = the RO's work** — competency rating, status, remarks, action, recognition — everything on one screen so a five-person queue moves fast.

Competency assessment (WCI022–025, BR001/007)

- **Technical + Behavioural** rated on the **Ranking Master's 5 levels** (1–5 chips here — real deployment shows the admin's ranking *titles*, ordered high→low). Colour reflects the value, not decoration.
- **Optional competencies handled** — "Automation Practices" shown dashed/un-rated: a department can configure fewer than five, and the un-scored state reads clearly.

Previous-week thread (WCI021/031, BR009)

- Last week's status/remarks/confidence sit in the evidence pane, with the **open action item and a Completed / Still-pending toggle** — closing the loop without touching *this* week's record.


Submission (WCI026–032, BR005/006)

- Status (On Track / Watch / Intervention), Confidence (H/M/L), **RO Remarks with a live 100-char counter**, Action Type, and **Action Notes that become required (min 100) whenever Action Type ≠ None** — the field would be disabled + optional at "None" per the exception table.
- **Recognition** (None / Appreciate / Appreciate Publicly) at the foot — acknowledgement, framed as the PO intended, not a reward.
- Footer states the **immutability + one-per-week** rule (BR003/008) right where Submit lives.


Edge cases baked in: no operational data → evidence pane shows "No operational data for this period" but review still allowed; no competency template → submission blocked with an admin nudge; already-submitted → view-only.

One deliberate contrast worth flagging: the queue pulls **direct reports only** (BR002) — one level — unlike the recursive own+direct+indirect visibility we set on Access Rights.

The retrospective surface — where "continuous performance history for PIE" stops being a promise and becomes something you can *see*. The design job here is **time**: the workspace captured one week; this shows the arc across many. Its distinctive element is a **competency-over-time matrix** — you read a person's trajectory in one glance, not by opening twelve reviews. Per-employee history: status arc, the competency trend grid, and the action-item track. Click any week to open its frozen snapshot.



And the **immutable snapshot** — what opening a past week gives you: the record exactly as submitted, evidence frozen at that moment, no edit affordances:



Weekly History — spec

The time lens (what makes this surface distinct)

- **Competency history matrix (WCI035)** is the centrepiece — competencies as rows, weeks as columns, each cell the frozen 1–5 rating in its Ranking-Master colour. You read trajectory instantly: Priya's Automation climbing 2→3 after training, Ownership dipping mid-January (the Watch trigger), Defect Detection steady at 5. This is the raw material the FRD keeps earmarking for **PIE (WCI010/BR010)**.
- **Status arc** across the 8 weeks — the green→amber transition makes the "when did this become a Watch" question a glance, not a hunt.
- **Action-item track** — every action raised, carried, or closed, with outcomes tied back (training completed → rating moved). This is the coaching-continuity story the module exists to build.

Immutability (WCI033/034, BR008)

- Opening any week gives the **frozen snapshot** — status, confidence, ratings, remarks, action, recognition, exactly as submitted, timestamped with who submitted.
- **Evidence is preserved as-captured, not recalculated** — the panel says so explicitly. This is the subtle-but-critical WCI034 point: if you re-derived Planning Accuracy from today's data, history would rewrite itself. The snapshot holds the numbers *as they were the day the review was written*.
- No edit affordances anywhere; footer states edits need admin override (BR008).

Manager history (WCI036) — the header carries "reviewed by" and each snapshot its submitter, so the complete RO review trail is preserved even across the reviewer-change edge case we noted (the RO at submission owns that week's record).